
CS2030 Lecture 11

Asynchronous Programming

Henry Chia (hchia@comp.nus.edu.sg)

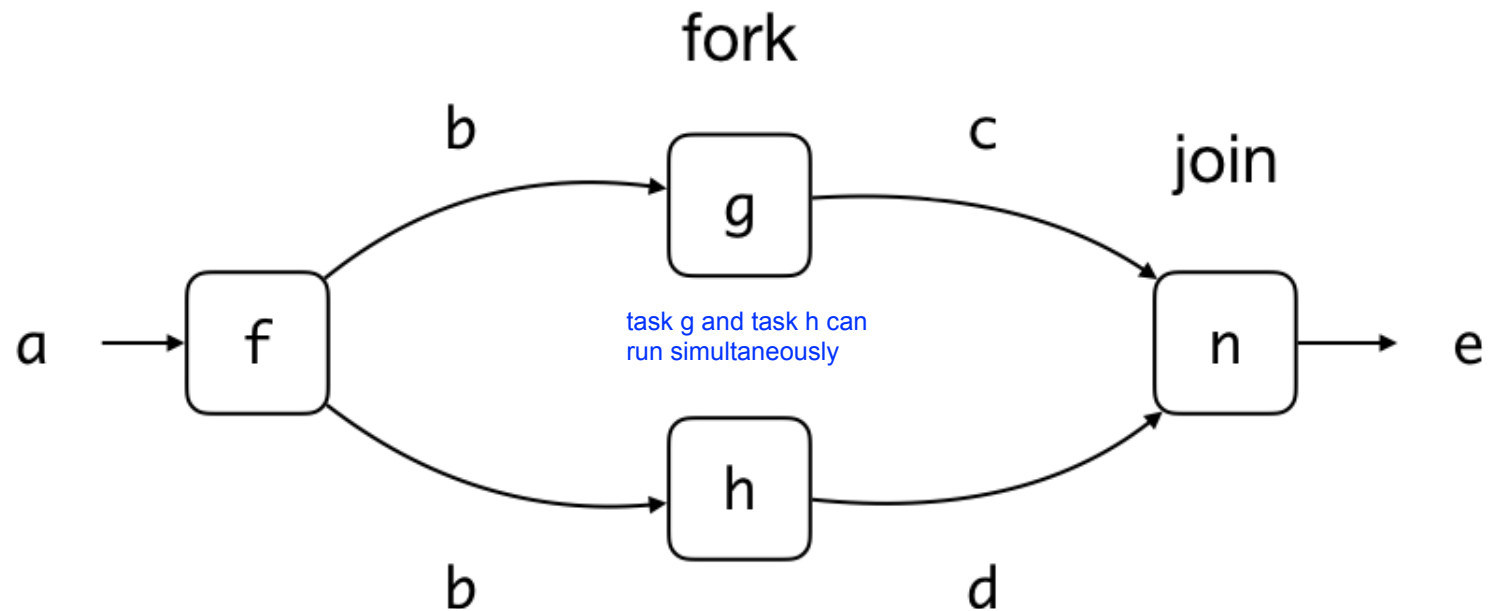
Semester 2 2022 / 2023

Lecture Outline

- Able to identify fork and join processes from a given computation graph
- Understand the difference between synchronous and asynchronous programming
- Appreciate asynchronous programming in the context of spawning threads to perform tasks
- Able to define asynchronous computations via Java's `CompletableFuture`
- Use of a callback to execute a block of code when an asynchronous task completes
- Able to convert synchronous code to an asynchronous version

Fork and Join

```
b = f(a);  
c = g(b);  
d = h(b);  
e = n(c,d);
```



- $f(a)$ invoked before $g(b)$ and $h(b)$; $n(c,d)$ invoked after
- If g and h does not produce side effects (i.e. does not depend or change external states), then
 - **fork** task g to execute at the same time as h , then
 - **join** back task g later

Synchronous Programming

```
int sleep(int n) { // to simulate a heavy task
    try {
        Thread.sleep(n * 1000); // Thread.sleep throws InterruptedException
    } catch (InterruptedException e) { }
    return n;
}
```

```
B f(A a) {
    System.out.println("f: start");
    sleep(5);
    System.out.println("f: done");
    return new B();
}
```

```
C g(B b, int n) {
    System.out.println("g: start");
    sleep(n);
    System.out.println("g: done");
    return new C();
}
```

```
D h(B b, int n) {
    System.out.println("h: start");
    sleep(n);
    System.out.println("h: done");
    return new D();
}
```

```
E n(C c, D d) {
    System.out.println("n: proceeds");
    return new E();
}
```

Synchronous Programming

- Synchronous programming: one function executes at a time

```
jshell> void foo(int m, int n) {  
    ...>     B b = f(new A());  
    ...>     C c = g(b, m);  
    ...>     D d = h(b, n);  
    ...>     E e = n(c, d);  
    ...> }  
| created method foo(int,int)  
  
jshell> foo(5, 10)  
f: start // f starts @ t = 0s  
f: done  // f completes @ t = 5s  
g: start // g starts after f completes @ t = 5s  
g: done  // g completes after five seconds @ t = 10s  
h: start // h starts after g completes @ t = 10s  
h: done  // h completes after another 10 seconds @ t = 20s  
n: proceeds // n proceeds @ t = 20s
```

- Since the execution of g and h can start at the same time
 - should require only 10 seconds to complete the execution of both methods, i.e. total time is $5 + 10 = 15$ seconds

Asynchronous Programming with Threads

- Spawn a separate process thread to compute g

```
jshell> void foo(int m, int n) throws InterruptedException {  
...>     B b = f(new A());  
...>     Thread t = new Thread(() -> g(b, m));  
...>     t.start();  
...>     h(b, n);  
...>     t.join(); // join() throws InterruptedException  
...>     System.out.println("n: proceeds");  
...> }  
| created method foo(int,int)
```

- A Runnable is passed to the Thread constructor
 - Runnable has the single abstract method void run()

```
jshell> void f() {}  
| created method f()
```

```
jshell> Runnable r = () -> f()  
r ==> $Lambda$...
```

```
jshell> r.run()
```

```
jshell> int g() { return 1;}  
| created method g()
```

```
jshell> r = () -> g()  
r ==> $Lambda$...
```

```
jshell> r.run()
```

Thread Completion via `join()`

- Wait for thread to complete using the `join()` method
 - `join()` method is blocking and returns when execution of the thread completes

```
jshell> foo(5, 10) // completes after 15 seconds
```

```
f: start
```

```
f: done
```

```
h: start
```

```
g: start
```

```
g: done
```

```
h: done // t.join() returns immediately as g has already completed
```

```
n: proceeds
```

```
jshell> foo(10, 5) // completes after 15 seconds
```

```
f: start
```

```
f: done
```

```
h: start
```

```
g: start
```

```
h: done
```

```
g: done // t.join() waits another 5 seconds for g to complete
```

```
n: proceeds
```

Java's `CompletableFuture<T>`

- Static methods `supplyAsync` (and `runAsync`) creates instances of `CompletableFuture` out of `Suppliers` (and `Runnables`)
 - returns a `CompletableFuture` rightaway
 - encapsulates the thread which starts execution

```
jshell> CompletableFuture.supplyAsync(() -> f(new A()))
$.. ==> java.util.concurrent.CompletableFuture@5d099f62[Not completed]
f: start

jshell> f: done
```

- Result of asynchronous computation is obtained via `join()`

```
jshell> CompletableFuture.supplyAsync(() -> f(new A())).join()
f: start
f: done
$.. ==> B@37bba400
```


Callback

- A *callback* is any executable code that is passed as an argument to other code
 - so that the former can be called back (or called after) upon completion of the latter
- Based on the *Hollywood Principle*:

“Don’t call us, we’ll call you (back)”
- A callback function is passed into **thenApply** (just a **map!**)

```
jshell> CompletableFuture.supplyAsync(() -> f(new A())).  
...> thenApply(x -> g(x, 5)).  
...> join()  
f: start  
f: done  
g: start  
g: done  
$.. ==> C@31cefde0
```

Callbacks in CompletableStage

- While `CompletableFuture` provides the static constructors, `CompletionStage` provides the other callback methods, e.g.
 - `thenAccept(Consumer<? super T> action)`
 - `thenApply(Function<? super T, ? extends U> fn)`
 - `thenCompose(Function<? super T, ? extends CompletableStage<U>> fn)`
 - `thenCombine(CompletionStage<? extends U> other, BiFunction<? super T, ? super U, ? extends V> fn)`
- `thenApply` and `thenCompose` are analogous to `map` and `flatMap` in `Optional`, `Stream`, etc.
- `CompletableFuture` is a `Functor` as well as a `Monad`!
- The `join()` method is blocking and returns the result when execution completes; returns `Void` for `Runnable` tasks
- ***Reminder:** Always ensure `that join() is called!`*

CompletableFuture Asynchronous Computation

- Constructing the CompletableFuture pipeline:

```
E foo(int m, int n) {  
    Supplier<B> suppB = () -> f(new A());  
  
    CompletableFuture<B> cfB = CompletableFuture.supplyAsync(suppB);  
    CompletableFuture<C> cfC = cfB.thenApply(x -> g(x, m));  
    CompletableFuture<D> cfD = cfB.thenApplyAsync(x -> h(x, n));  
    CompletableFuture<E> cfE = cfC.thenCombine(cfD, (c,d) -> n(c, d));  
  
    E e = cfE.join();  
  
    return e;  
}
```

```
jshell> foo(5, 10)  
f: start // t = 0s  
f: done  // t = 5s  
g: start // t = 5s  
h: start // t = 5s  
g: done  // t = 10s  
h: done  // t = 15s  
$.. ==> E@49097b5d
```

```
jshell> foo(10, 5)  
f: start // t = 0s  
f: done  // t = 5s  
g: start // t = 5s  
h: start // t = 5s  
h: done  // t = 10s  
g: done  // t = 15s  
$.. ==> E@37a71e93
```

Converting Synchronous to Asynchronous

- Give the following synchronous program fragment

```
int foo(int x) {  
    if (x < 0) {  
        return 0;  
    } else {  
        return doWork(x);  
    }  
}
```

- The asynchronous version is

```
CompletableFuture<Integer> fooAsync(int x) {  
    if (x < 0) {  
        return CompletableFuture.completedFuture(0);  
    } else {  
        return CompletableFuture.supplyAsync(() -> doWork(x));  
    }  
}
```

Completed Value

- `CompletableFuture.completedFuture(U value)` wraps a completed value in a `CompletableFuture`

Converting Synchronous to Asynchronous

- Suppose we have the following synchronous method calls

```
int y = foo(5)
int z = bar(y)
```

with bar defined as

```
int bar(int x) {
    return x;
}
```

- The above sequence of function calls can be composed as

```
int z = bar(foo(5))
```

- The equivalent asynchronous version is

```
int z = fooAsync(5).
    thenApply(x -> bar(x)).
    join();
```

Converting Synchronous to Asynchronous

- What if we switch the method calls, i.e.

```
int y = bar(5)
int z = foo(y)
```

and suppose bar is asynchronous as well, i.e.

```
CompletableFuture<Integer> barAsync(int x) {
    return CompletableFuture.completedFuture(x);
}
```

- Then the equivalent asynchronous version is

```
int z = barAsync(5).
    thenCompose(y -> fooAsync(y)).
    join();
```

barAsync also returns a CompletableFuture
must use thenCompose (flatMap equivalent)

- What if we use thenApply instead of thenCompose?

cannot compile or will have nested contexts(CompletableFuture)

Combining Completable Futures

- Combine results of two CompletableFutures via BiFunction

```
int z = fooAsync(5).  
    thenCombine(barAsync(5), (x,y) -> x + y).  
    join()
```

- Both fooAsync and barAsync must be completed, before resulting CompletableFuture from thenCombine completes
- To summarize...
 - use runAsync or supplyAsync to create
 - then<X><Y><Z> where
 - ▷ X is Accept, Combine, Compose, Run, ...
 - ▷ Y is nothing, Both, Either
 - ▷ Z is nothing or Async

Async Variants of Callback Methods

- Callback methods have an Async variant, e.g.
 - `thenRun/thenRunAsync` *may* run on a different thread
 - `thenRun` runs on the same thread if still busy

```
jshell> void foo() {  
...>     CompletableFuture<Void> cf1 = CompletableFuture.runAsync(() -> {  
...>         sleep(5);  
...>         System.out.println("cf1: " + Thread.currentThread().getName()); });  
...>  
...>     CompletableFuture<Void> cf2 = cf1.thenRun(() -> {  
...>         sleep(5);  
...>         System.out.println("cf2: " + Thread.currentThread().getName()); });  
...>  
...>     CompletableFuture<Void> cf3 = cf1.thenRunAsync(() -> {  
...>         sleep(5);  
...>         System.out.println("cf3: " + Thread.currentThread().getName()); });  
...>  
...>     cf2.join();  
...>     cf3.join();  
...> }  
| created method foo()  
  
jshell> foo()  
cf1: ForkJoinPool.commonPool-worker-3  
cf2: ForkJoinPool.commonPool-worker-3  
cf3: ForkJoinPool.commonPool-worker-5
```